

# MaunaLoa

April 1, 2025

## 1 Mauna Loa Daily Temps Time Series

**References:** - ARIMA Model in Python Time Series Forecasting #6. Nachiketa Hebbar - ARIMA for Time Series Forecasting; A Complete Guide. Zaina Saadeddin - Time Series Forecasting with ARIMA: Everything You Need to Know!. Nayeem Islam

```
[ ]: # Install pmdarima library  
# pmdarima does not support Numpy 2.0  
# faced challenges trying to get pmdarima to install  
# repeated uninstalls and installs of numpy and pandas  
# decided not to build a user interface due to the challenges with pmdarima
```

```
[ ]: !python --version
```

Python 3.11.11

- python 3.11.11
- numpy 1.23.2
- pandas 2.2.2

```
[ ]: !pip uninstall numpy
```

Found existing installation: numpy 2.2.4

Uninstalling numpy-2.2.4:

Would remove:

```
/usr/local/bin/f2py  
/usr/local/bin/numpy-config  
/usr/local/lib/python3.11/dist-packages/numpy-2.2.4.dist-info/*  
/usr/local/lib/python3.11/dist-  
packages/numpy.libs/libgfortran-040039e1-0352e75f.so.5.0.0  
/usr/local/lib/python3.11/dist-  
packages/numpy.libs/libquadmath-96973f99-934c22de.so.0.0.0  
/usr/local/lib/python3.11/dist-  
packages/numpy.libs/libscipy_openblas64_-6bb31eeb.so  
/usr/local/lib/python3.11/dist-packages/numpy/*
```

Proceed (Y/n)? y

Successfully uninstalled numpy-2.2.4

```
[ ]: !pip uninstall pandas
```

```
Found existing installation: pandas 2.2.2
Uninstalling pandas-2.2.2:
  Would remove:
    /usr/local/lib/python3.11/dist-packages/pandas-2.2.2.dist-info/*
    /usr/local/lib/python3.11/dist-packages/pandas/*
Proceed (Y/n)? y
  Successfully uninstalled pandas-2.2.2
```

```
[ ]: !pip uninstall pmdarima
```

```
Found existing installation: pmdarima 2.0.4
Uninstalling pmdarima-2.0.4:
  Would remove:
    /usr/local/lib/python3.11/dist-packages/pmdarima-2.0.4.dist-info/*
    /usr/local/lib/python3.11/dist-packages/pmdarima/*
Proceed (Y/n)? y
  Successfully uninstalled pmdarima-2.0.4
```

```
[ ]: !pip install numpy==1.26.4
```

```
Collecting numpy==1.26.4
  Downloading
numpy-1.26.4-cp311-cp311-manylinux_2_17_x86_64.manylinux2014_x86_64.whl.metadata
(61 kB)
                                0.0/61.0
kB ? eta -:--:--
                                61.0/61.0 kB 3.4
MB/s eta 0:00:00
Downloading
numpy-1.26.4-cp311-cp311-manylinux_2_17_x86_64.manylinux2014_x86_64.whl (18.3
MB)
                                18.3/18.3 MB
36.4 MB/s eta 0:00:00
Installing collected packages: numpy
  Attempting uninstall: numpy
    Found existing installation: numpy 2.2.4
    Uninstalling numpy-2.2.4:
      Successfully uninstalled numpy-2.2.4
  Successfully installed numpy-1.26.4
```

```
[ ]: !pip install pandas==2.2.2
```

```
Requirement already satisfied: pandas==2.2.2 in /usr/local/lib/python3.11/dist-
packages (2.2.2)
Requirement already satisfied: numpy>=1.23.2 in /usr/local/lib/python3.11/dist-
packages (from pandas==2.2.2) (1.26.4)
Requirement already satisfied: python-dateutil>=2.8.2 in
```

```
/usr/local/lib/python3.11/dist-packages (from pandas==2.2.2) (2.8.2)
Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.11/dist-
packages (from pandas==2.2.2) (2025.2)
Requirement already satisfied: tzdata>=2022.7 in /usr/local/lib/python3.11/dist-
packages (from pandas==2.2.2) (2025.2)
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.11/dist-
packages (from python-dateutil>=2.8.2->pandas==2.2.2) (1.17.0)
```

```
[ ]: !pip install pmdarima
```

```
Collecting pmdarima
  Using cached pmdarima-2.0.4-cp311-cp311-
manylinux_2_17_x86_64.manylinux2014_x86_64.manylinux_2_28_x86_64.whl.metadata
(7.8 kB)
Requirement already satisfied: joblib>=0.11 in /usr/local/lib/python3.11/dist-
packages (from pmdarima) (1.4.2)
Requirement already satisfied: Cython!=0.29.18,!=0.29.31,>=0.29 in
/usr/local/lib/python3.11/dist-packages (from pmdarima) (3.0.12)
Requirement already satisfied: numpy>=1.21.2 in /usr/local/lib/python3.11/dist-
packages (from pmdarima) (1.26.4)
Requirement already satisfied: pandas>=0.19 in /usr/local/lib/python3.11/dist-
packages (from pmdarima) (2.2.2)
Requirement already satisfied: scikit-learn>=0.22 in
/usr/local/lib/python3.11/dist-packages (from pmdarima) (1.6.1)
Requirement already satisfied: scipy>=1.3.2 in /usr/local/lib/python3.11/dist-
packages (from pmdarima) (1.14.1)
Requirement already satisfied: statsmodels>=0.13.2 in
/usr/local/lib/python3.11/dist-packages (from pmdarima) (0.14.4)
Requirement already satisfied: urllib3 in /usr/local/lib/python3.11/dist-
packages (from pmdarima) (2.3.0)
Requirement already satisfied: setuptools!=50.0.0,>=38.6.0 in
/usr/local/lib/python3.11/dist-packages (from pmdarima) (75.2.0)
Requirement already satisfied: packaging>=17.1 in
/usr/local/lib/python3.11/dist-packages (from pmdarima) (24.2)
Requirement already satisfied: python-dateutil>=2.8.2 in
/usr/local/lib/python3.11/dist-packages (from pandas>=0.19->pmdarima) (2.8.2)
Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.11/dist-
packages (from pandas>=0.19->pmdarima) (2025.2)
Requirement already satisfied: tzdata>=2022.7 in /usr/local/lib/python3.11/dist-
packages (from pandas>=0.19->pmdarima) (2025.2)
Requirement already satisfied: threadpoolctl>=3.1.0 in
/usr/local/lib/python3.11/dist-packages (from scikit-learn>=0.22->pmdarima)
(3.6.0)
Requirement already satisfied: patsy>=0.5.6 in /usr/local/lib/python3.11/dist-
packages (from statsmodels>=0.13.2->pmdarima) (1.0.1)
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.11/dist-
packages (from python-dateutil>=2.8.2->pandas>=0.19->pmdarima) (1.17.0)
Using cached pmdarima-2.0.4-cp311-cp311-
```

manylinux\_2\_17\_x86\_64.manylinux2014\_x86\_64.manylinux\_2\_28\_x86\_64.whl (2.2 MB)  
Installing collected packages: pmdarima  
Successfully installed pmdarima-2.0.4

```
[ ]: # Load libraries
import pandas as pd
import numpy as np
import pmdarima as pm
```

## 1.1 Read Data: MaunaLoa Daily Temperatures

```
[ ]: # read csv file
df = pd.read_csv('/content/MaunaLoaDailyTemps-1.csv', index_col='DATE',
                parse_dates=True)

# drop missing values
df = df.dropna()

# Show dataset
print('Shape of data', df.shape)
df.head()
```

Shape of data (1821, 5)


```
[ ]:
```

	MinTemp	MaxTemp	AvgTemp	Sunrise	Sunset
DATE					
2014-01-01	33.0	46.0	40.0	657	1756
2014-01-02	35.0	50.0	43.0	657	1756
2014-01-03	36.0	45.0	41.0	657	1757
2014-01-04	32.0	41.0	37.0	658	1757
2014-01-05	24.0	38.0	31.0	658	1758

The dataset has 1,821 rows and 5 columns.

```
[ ]: # plot data
df['AvgTemp'].plot(figsize=(12,5))
```

```
[ ]: <Axes: xlabel='DATE'>
```



## 1.2 Is Data Stationary?

If P-value  $< 0.05$ , data is stationary

If P-value  $> 0.05$ , data is not stationary. Data has an increasing or decreasing trend/p>

```
[ ]: # Function to check whether data is stationary or not
# Modeled time series data needs to be stationary
# The time series mean, variance, etc are constant over time

from statsmodels.tsa.stattools import adfuller

def ad_test(dataset):
    dfctest = adfuller(dataset, autolag = 'AIC')
    print("1. ADF: ", dfctest[0])
    print("2. P-Value: ", dfctest[1])
    print("3. Num of Lags: ", dfctest[2])
    print("4. Num of Observations Used for ADF Regression and Critical Values_↵
    ↵Calculation: ", dfctest[3])
    print("5. Critical Values : ")
    for key, val in dfctest[4].items():
        print("\t", key, ": ", val)

    # Interpret the results
    if dfctest[1] > 0.05:
        print("The data is not stationary.")
    else:
        print("The data is stationary.")
```

```
[ ]: # P-value should be as low as possible. < 0.05
```

```
ad_test(df['AvgTemp'])
```

```
1. ADF: -6.554680125068777
2. P-Value: 8.675937480199653e-09
3. Num of Lags: 12
4. Num of Observations Used for ADF Regression and Critical Values Calculation:
1808
5. Critical Values :
    1% : -3.433972018026501
    5% : -2.8631399192826676
   10% : -2.5676217442756872
The data is stationary.
```

### 1.3 Find the Best ARIMA Model

```
[ ]: # Load auto_arima
from pmdarima import auto_arima
# ignore warnings
import warnings
warnings.filterwarnings("ignore")
```

```
[ ]: # Find the best ARIMA model
stepwise_fit = auto_arima(df['AvgTemp'], trace=True, suppress_warnings=True)

stepwise_fit.summary()
```

Performing stepwise search to minimize aic

```
ARIMA(2,0,2)(0,0,0)[0] intercept : AIC=8344.735, Time=3.55 sec
ARIMA(0,0,0)(0,0,0)[0] intercept : AIC=10347.755, Time=0.07 sec
ARIMA(1,0,0)(0,0,0)[0] intercept : AIC=8365.701, Time=0.21 sec
ARIMA(0,0,1)(0,0,0)[0] intercept : AIC=9136.225, Time=0.43 sec
ARIMA(0,0,0)(0,0,0)[0]          : AIC=19192.139, Time=0.05 sec
ARIMA(1,0,2)(0,0,0)[0] intercept : AIC=8355.947, Time=2.80 sec
ARIMA(2,0,1)(0,0,0)[0] intercept : AIC=8356.308, Time=4.99 sec
ARIMA(3,0,2)(0,0,0)[0] intercept : AIC=8347.311, Time=4.13 sec
ARIMA(2,0,3)(0,0,0)[0] intercept : AIC=8318.337, Time=4.95 sec
ARIMA(1,0,3)(0,0,0)[0] intercept : AIC=8330.192, Time=5.30 sec
ARIMA(3,0,3)(0,0,0)[0] intercept : AIC=8310.577, Time=4.87 sec
ARIMA(4,0,3)(0,0,0)[0] intercept : AIC=8332.293, Time=7.24 sec
ARIMA(3,0,4)(0,0,0)[0] intercept : AIC=8317.620, Time=5.89 sec
ARIMA(2,0,4)(0,0,0)[0] intercept : AIC=8306.228, Time=7.34 sec
ARIMA(1,0,4)(0,0,0)[0] intercept : AIC=8297.028, Time=4.73 sec
ARIMA(0,0,4)(0,0,0)[0] intercept : AIC=8455.435, Time=1.50 sec
ARIMA(1,0,5)(0,0,0)[0] intercept : AIC=8295.034, Time=9.00 sec
ARIMA(0,0,5)(0,0,0)[0] intercept : AIC=8419.091, Time=1.62 sec
ARIMA(2,0,5)(0,0,0)[0] intercept : AIC=8302.544, Time=8.08 sec
ARIMA(1,0,5)(0,0,0)[0]          : AIC=8304.533, Time=0.58 sec
```

Best model: ARIMA(1,0,5)(0,0,0)[0] intercept  
Total fit time: 77.383 seconds

[ ]:

<b>Dep. Variable:</b>	y	<b>No. Observations:</b>	1821
<b>Model:</b>	SARIMAX(1, 0, 5)	<b>Log Likelihood</b>	-4139.517
<b>Date:</b>	Mon, 31 Mar 2025	<b>AIC</b>	8295.034
<b>Time:</b>	21:09:48	<b>BIC</b>	8339.092
<b>Sample:</b>	0	<b>HQIC</b>	8311.288
	- 1821		
<b>Covariance Type:</b>	opg		

	coef	std err	z	P>  z	[0.025	0.975]
<b>intercept</b>	1.2070	0.362	3.335	0.001	0.498	1.916
<b>ar.L1</b>	0.9739	0.008	124.621	0.000	0.959	0.989
<b>ma.L1</b>	-0.1246	0.024	-5.256	0.000	-0.171	-0.078
<b>ma.L2</b>	-0.2196	0.024	-9.115	0.000	-0.267	-0.172
<b>ma.L3</b>	-0.2056	0.024	-8.615	0.000	-0.252	-0.159
<b>ma.L4</b>	-0.1373	0.023	-6.034	0.000	-0.182	-0.093
<b>ma.L5</b>	-0.0476	0.024	-1.952	0.051	-0.095	0.000
<b>sigma2</b>	5.4981	0.172	31.952	0.000	5.161	5.835

<b>Ljung-Box (L1) (Q):</b>	0.00	<b>Jarque-Bera (JB):</b>	20.20
<b>Prob(Q):</b>	0.95	<b>Prob(JB):</b>	0.00
<b>Heteroskedasticity (H):</b>	0.81	<b>Skew:</b>	-0.17
<b>Prob(H) (two-sided):</b>	0.01	<b>Kurtosis:</b>	3.39

Warnings:

[1] Covariance matrix calculated using the outer product of gradients (complex-step).

The best arima model is ARIMA(1,0,5). - p: The number of past values (lags) considered in the AR term. 1 - d: The degree of differencing applied to the data. 0 - q: The number of past forecast errors included in the MA term. 5

```
[ ]: # Load ARIMA
from statsmodels.tsa.arima.model import ARIMA
```

## 1.4 Split Data into Train and Test

```
[ ]: print(df.shape)
train = df.iloc[:-30] # all values except the last 30 values
test = df.iloc[-30:] #last 30 values

print(train.shape, test.shape)
```

(1821, 5)  
(1791, 5) (30, 5)

## 1.5 Train the Model

```
[ ]: model = ARIMA(train['AvgTemp'], order=(1,0,5))
model = model.fit()
model.summary()
```

```
[ ]:
```

<b>Dep. Variable:</b>	AvgTemp	<b>No. Observations:</b>	1791
<b>Model:</b>	ARIMA(1, 0, 5)	<b>Log Likelihood</b>	-4070.198
<b>Date:</b>	Mon, 31 Mar 2025	<b>AIC</b>	8156.395
<b>Time:</b>	21:10:04	<b>BIC</b>	8200.320
<b>Sample:</b>	0	<b>HQIC</b>	8172.614
	- 1791		
<b>Covariance Type:</b>	opg		

	coef	std err	z	P>  z	[0.025	0.975]
<b>const</b>	46.5856	0.758	61.454	0.000	45.100	48.071
<b>ar.L1</b>	0.9856	0.005	188.230	0.000	0.975	0.996
<b>ma.L1</b>	-0.1412	0.023	-6.124	0.000	-0.186	-0.096
<b>ma.L2</b>	-0.2268	0.024	-9.635	0.000	-0.273	-0.181
<b>ma.L3</b>	-0.2168	0.023	-9.251	0.000	-0.263	-0.171
<b>ma.L4</b>	-0.1479	0.023	-6.491	0.000	-0.193	-0.103
<b>ma.L5</b>	-0.0595	0.024	-2.438	0.015	-0.107	-0.012
<b>sigma2</b>	5.5093	0.174	31.624	0.000	5.168	5.851

<b>Ljung-Box (L1) (Q):</b>	0.00	<b>Jarque-Bera (JB):</b>	14.88
<b>Prob(Q):</b>	0.97	<b>Prob(JB):</b>	0.00
<b>Heteroskedasticity (H):</b>	0.82	<b>Skew:</b>	-0.15
<b>Prob(H) (two-sided):</b>	0.01	<b>Kurtosis:</b>	3.33

Warnings:

[1] Covariance matrix calculated using the outer product of gradients (complex-step).

```
[ ]: # Make a prediction on the test data then compare to actual

start=len(train)
end=len(train)+len(test)-1
#pred=model.predict(start=start,end=end,typ='levels').rename('ARIMA_
↳Predictions').rename('ARIMA Predictions')
pred=model.predict(start=start,end=end,typ='levels').rename('ARIMA Predictions')
pred.index = df.index[start:end+1]
print(pred)
```

DATE

```
2018-12-01    44.754109
2018-12-02    44.987795
2018-12-03    45.388741
2018-12-04    45.721545
2018-12-05    45.863733
2018-12-06    45.874126
```



2018-12-07	45.884370
2018-12-08	45.894466
2018-12-09	45.904417
2018-12-10	45.914225
2018-12-11	45.923891
2018-12-12	45.933418
2018-12-13	45.942808
2018-12-14	45.952063
2018-12-15	45.961185
2018-12-16	45.970175
2018-12-17	45.979036
2018-12-18	45.987769
2018-12-19	45.996377
2018-12-20	46.004861
2018-12-21	46.013222
2018-12-22	46.021463
2018-12-23	46.029586
2018-12-24	46.037591
2018-12-25	46.045481
2018-12-26	46.053258
2018-12-27	46.060923
2018-12-28	46.068477
2018-12-29	46.075922
2018-12-30	46.083261

Name: ARIMA Predictions, dtype: float64

```
[ ]: print(test['AvgTemp'])
```

DATE	
2018-12-01	44.0
2018-12-02	42.0
2018-12-03	45.0
2018-12-04	48.0
2018-12-05	45.0
2018-12-06	44.0
2018-12-07	45.0
2018-12-08	44.0
2018-12-09	45.0
2018-12-10	46.0
2018-12-11	47.0
2018-12-12	47.0
2018-12-13	45.0
2018-12-14	46.0
2018-12-15	47.0
2018-12-16	49.0
2018-12-17	51.0
2018-12-18	43.0
2018-12-19	42.0

```

2018-12-20    48.0
2018-12-21    50.0
2018-12-22    47.0
2018-12-23    47.0
2018-12-24    44.0
2018-12-25    42.0
2018-12-26    40.0
2018-12-27    39.0
2018-12-28    40.0
2018-12-29    42.0
2018-12-30    46.0
Name: AvgTemp, dtype: float64

```

```


[ ]: pred.plot(legend=True)
test['AvgTemp'].plot(legend=True)

```

```

[ ]: <Axes: xlabel='DATE'>


```



## 1.6 Plot Train, Test, & Predicted Data

```
[ ]: # Load matplotlib
import matplotlib.pyplot as plt

# Define plot
plt.figure(figsize=(14,7))
plt.plot(train.index, train["AvgTemp"], label='Train', color='#203147') # train_
    ↳data
plt.plot(test.index, test["AvgTemp"], label='Test', color='#01ef63') # test_
    ↳data
plt.plot(test.index, pred, label='Prediction', color='red') #
    ↳forecasted data
plt.title('Temperature: Train, Test, Predicted')
plt.xlabel('Date')
plt.ylabel('Temperature')
plt.legend()
plt.show()
```



```
[32]: test['AvgTemp'].mean()
```

```
[32]: 45.0
```

The average for the test data is 45.

```
[ ]: pred = pred[:len(test)]
test_temp = test['AvgTemp'][:len(pred)]
```

```
# Calculate RMSE

from sklearn.metrics import mean_squared_error
rmse = np.sqrt(mean_squared_error(test_temp,pred))
print(f"RMSE: {rmse: .4f}")
```

RMSE: 3.0005

On average predictions are off by +/- 3 degrees.

## 1.7 Forecast Past the End of the Original Data

Forecast: 2018-12-31 to 2109-04-09

```
[ ]: # Train the model on the entire dataset
model2 = ARIMA(df['AvgTemp'], order=(1,0,5))

# Fit the model
model2 = model2.fit()

# Model Summary
model2.summary()
```

```
[ ]:
```

<b>Dep. Variable:</b>	AvgTemp	<b>No. Observations:</b>	1821
<b>Model:</b>	ARIMA(1, 0, 5)	<b>Log Likelihood</b>	-4138.130
<b>Date:</b>	Mon, 31 Mar 2025	<b>AIC</b>	8292.261
<b>Time:</b>	21:18:20	<b>BIC</b>	8336.318
<b>Sample:</b>	0	<b>HQIC</b>	8308.514
	- 1821		
<b>Covariance Type:</b>	opg		

	coef	std err	z	P>  z	[0.025	0.975]
const	46.5285	0.758	61.403	0.000	45.043	48.014
ar.L1	0.9860	0.005	192.909	0.000	0.976	0.996
ma.L1	-0.1403	0.023	-6.125	0.000	-0.185	-0.095
ma.L2	-0.2328	0.023	-10.038	0.000	-0.278	-0.187
ma.L3	-0.2163	0.023	-9.280	0.000	-0.262	-0.171
ma.L4	-0.1478	0.023	-6.542	0.000	-0.192	-0.104
ma.L5	-0.0587	0.024	-2.424	0.015	-0.106	-0.011
sigma2	5.5080	0.173	31.925	0.000	5.170	5.846

<b>Ljung-Box (L1) (Q):</b>	0.00	<b>Jarque-Bera (JB):</b>	15.69
<b>Prob(Q):</b>	0.98	<b>Prob(JB):</b>	0.00
<b>Heteroskedasticity (H):</b>	0.81	<b>Skew:</b>	-0.15
<b>Prob(H) (two-sided):</b>	0.01	<b>Kurtosis:</b>	3.34

Warnings:

[1] Covariance matrix calculated using the outer product of gradients (complex-step).


```
[ ]: # Forecast future values: 100 days
forecast = model2.forecast(steps = 100)

# assign dates as index
forecast.index = pd.date_range(start=df.index[-1], periods=101,
    ↪inclusive="right")

# print forecast
print(forecast)
```

```
2018-12-31    46.418166
2019-01-01    46.113912
2019-01-02    45.617874
2019-01-03    45.249566
2019-01-04    45.116916
...
2019-04-05    46.136883
2019-04-06    46.142362
2019-04-07    46.147764
2019-04-08    46.153091
2019-04-09    46.158343
Freq: D, Name: predicted_mean, Length: 100, dtype: float64
```

```
[ ]: # Define plot for forecast
plt.figure(figsize=(14,7))
plt.plot(forecast.index,forecast, label='Forecast', color='red')    # ↪
    ↪forecasted data
plt.title('Temperature: Forecast 100 days')
plt.xlabel('Date')
plt.ylabel('Temperature')
plt.legend()
plt.show()
```



## 1.8 Save Model to a File

```
[36]: import joblib
      joblib.dump(model2, "manualoa_arima_model_jl.sav")
```

```
[36]: ['manualoa_arima_model_jl.sav']
```

```
[37]: load_model = joblib.load("manualoa_arima_model_jl.sav")
```

```
[40]: temp_forecast=load_model.forecast(steps=10)
      last_date = pd.Timestamp("2018-12-30")
      days = 10
      temp_forecast.index =pd.date_range(last_date,periods = days + 1, freq="D")[1:]
      print(temp_forecast)
```

```
2018-12-31    46.418166
2019-01-01    46.113912
2019-01-02    45.617874
2019-01-03    45.249566
2019-01-04    45.116916
2019-01-05    45.136666
2019-01-06    45.156140
2019-01-07    45.175341
2019-01-08    45.194274
2019-01-09    45.212942
Freq: D, Name: predicted_mean, dtype: float64
```