

Outliers

February 15, 2024

1 Finding Outliers

Outliers are extreme values in a dataset. Are these true values or the result of an error?

Below is a sampling of methods to find outliers

I normally use boxplot.

```
[130]: import matplotlib.pyplot as plt
import numpy as np
import pandas as pd
import seaborn as sns
```

1.1 Load Diabetes dataset

```
[132]: # Convert sklearn diabetes dataset to dataframe
def sklearn_to_df(sklearn_dataset):
    df = pd.DataFrame(sklearn_dataset.data, columns=sklearn_dataset.
        ↪feature_names)
    df['target'] = pd.Series(sklearn_dataset.target)
    return df

df_diabetes = sklearn_to_df(datasets.load_diabetes())
```

```
[133]: # Diabetes dataset dimensions
# 442 rows and 11 columns

df_diabetes.shape
```

```
[133]: (442, 11)
```

1.2 Method 1: Sorting using SQL

- Oracle LIVE SQL
- Sort column in ascending order.
- Look at first 10 rows.

There are no outliers.

| STORE_NAME | TOTAL_SALES |
|---------------|-------------|
| São Paulo | 3148.22 |
| Tokyo | 3263.82 |
| Buenos Aires | 3495.72 |
| New York City | 3582.33 |
| Perth | 3707.49 |
| Chicago | 3721.28 |
| Berlin | 3791.11 |
| Beijing | 3849.33 |
| Johannesburg | 3870.98 |
| Utrecht | 3934.56 |

- Oracle LIVE SQL
- Sort column in descending order.
- Look at first 10 rows.

The first two values are outliers.

| STORE_NAME | TOTAL_SALES |
|---------------|-------------|
| - | 299889.62 |
| Online | 211107.78 |
| New Dehli | 5291.76 |
| Sydney | 4605.82 |
| Tel Aviv | 4457.19 |
| Mumbai | 4443.81 |
| London | 4427.49 |
| San Francisco | 4380.91 |
| Seattle | 4294.2 |
| Madrid | 4187.88 |

1.3 Method 2: Boxplot using Python

Visualize outliers using a boxplot

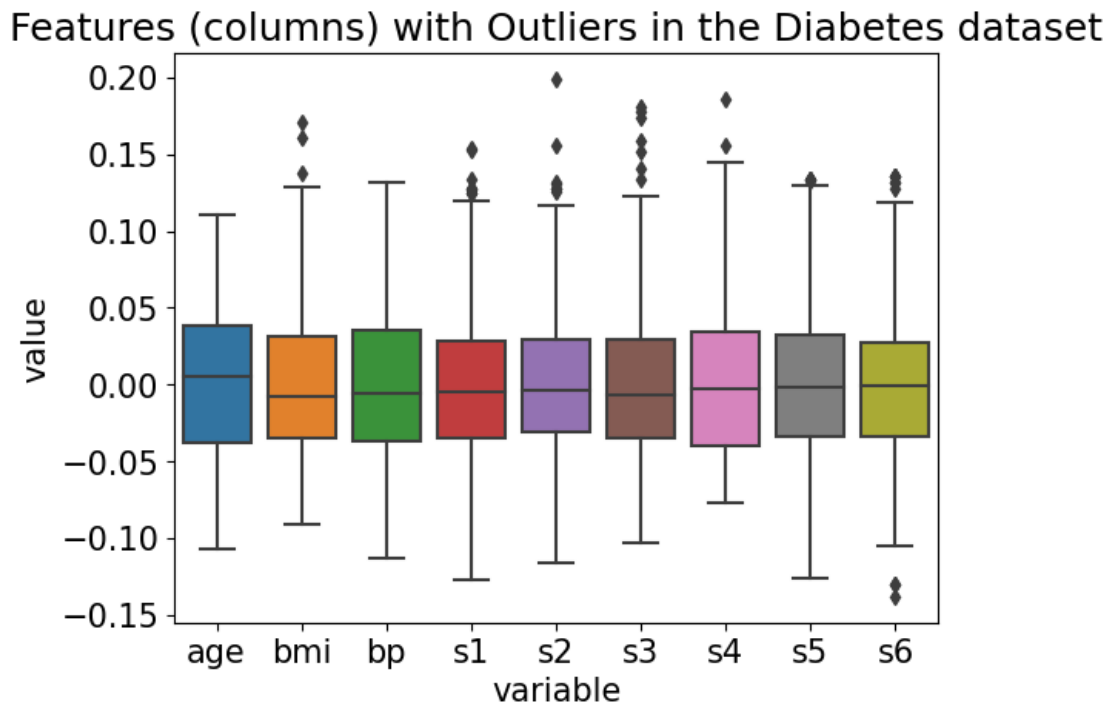
The markings above and below the box whiskers are outliers.

```
[140]: # Assign column names to a variable
cols = ['age', 'bmi', 'bp', 's1', 's2', 's3', 's4', 's5', 's6']
```

```
[141]: # Generate a boxplot for each column name in the cols variable.

df_diabetes_g = df_diabetes[cols]
sns.boxplot(x="variable", y = "value", data=pd.melt(df_diabetes_g))
plt.title("Features (columns) with Outliers in the Diabetes dataset")
```

```
[141]: Text(0.5, 1.0, 'Features (columns) with Outliers in the Diabetes dataset')
```



1.4 Method 3: Calculate Interquantile Range (IQR) using Python

A measure of the spread of the data

1.4.1 Use the describe() method to get the quantiles for each column

```
[144]: # Descriptive statistics for the Diabetes dataset
# First quantile (25%)
# Second quantile (50%)
# Third quantile (75%)

q = df_diabetes.describe().transpose()
q[['25%', '50%', '75%']]
```

```
[144]:
```

| | 25% | 50% | 75% |
|--------|-----------|------------|------------|
| age | -0.037299 | 0.005383 | 0.038076 |
| sex | -0.044642 | -0.044642 | 0.050680 |
| bmi | -0.034229 | -0.007284 | 0.031248 |
| bp | -0.036656 | -0.005670 | 0.035644 |
| s1 | -0.034248 | -0.004321 | 0.028358 |
| s2 | -0.030358 | -0.003819 | 0.029844 |
| s3 | -0.035117 | -0.006584 | 0.029312 |
| s4 | -0.039493 | -0.002592 | 0.034309 |
| s5 | -0.033246 | -0.001947 | 0.032432 |
| s6 | -0.033179 | -0.001078 | 0.027917 |
| target | 87.000000 | 140.500000 | 211.500000 |

1.4.2 Calculate the IQR for each column using the 25% and 75% quantiles

```
[146]: q['IQR'] = q['75%'] - q['25%']
q[['25%', '75%', 'IQR']]
```

```
[146]:
```

| | 25% | 75% | IQR |
|--------|-----------|------------|------------|
| age | -0.037299 | 0.038076 | 0.075375 |
| sex | -0.044642 | 0.050680 | 0.095322 |
| bmi | -0.034229 | 0.031248 | 0.065477 |
| bp | -0.036656 | 0.035644 | 0.072300 |
| s1 | -0.034248 | 0.028358 | 0.062606 |
| s2 | -0.030358 | 0.029844 | 0.060203 |
| s3 | -0.035117 | 0.029312 | 0.064429 |
| s4 | -0.039493 | 0.034309 | 0.073802 |
| s5 | -0.033246 | 0.032432 | 0.065678 |
| s6 | -0.033179 | 0.027917 | 0.061096 |
| target | 87.000000 | 211.500000 | 124.500000 |

1.4.3 Calculate the Upper and Lower limits for each column

```
[148]: q['Upper'] = q['75%'] + (1.5 * q['IQR'])
q['Lower'] = q['25%'] - (1.5 * q['IQR'])

q[['25%', '75%', 'IQR', 'Upper', 'Lower']]
```

```
[148]:
```

| | 25% | 75% | IQR | Upper | Lower |
|-----|-----------|----------|----------|----------|-----------|
| age | -0.037299 | 0.038076 | 0.075375 | 0.151139 | -0.150362 |
| sex | -0.044642 | 0.050680 | 0.095322 | 0.193663 | -0.187624 |
| bmi | -0.034229 | 0.031248 | 0.065477 | 0.129464 | -0.132445 |
| bp | -0.036656 | 0.035644 | 0.072300 | 0.144094 | -0.145106 |
| s1 | -0.034248 | 0.028358 | 0.062606 | 0.122267 | -0.128157 |
| s2 | -0.030358 | 0.029844 | 0.060203 | 0.120149 | -0.120663 |
| s3 | -0.035117 | 0.029312 | 0.064429 | 0.125954 | -0.131760 |

| | | | | | |
|--------|-----------|------------|------------|------------|------------|
| s4 | -0.039493 | 0.034309 | 0.073802 | 0.145012 | -0.150197 |
| s5 | -0.033246 | 0.032432 | 0.065678 | 0.130949 | -0.131762 |
| s6 | -0.033179 | 0.027917 | 0.061096 | 0.119561 | -0.124823 |
| target | 87.000000 | 211.500000 | 124.500000 | 398.250000 | -99.750000 |

1.4.4 List the outliers for a given column

The outliers are identified as having a value greater than the upper or less than the lower

```
[150]: # Upper Outliers for bmi
df_diabetes.loc[(df_diabetes['bmi']>0.13), 'bmi']
```

```
[150]: 256    0.160855
      366    0.137143
      367    0.170555
      Name: bmi, dtype: float64
```

```
[151]: # Upper Outliers for s1
df_diabetes.loc[(df_diabetes['s1']>0.12), 's1']
```

```
[151]: 123    0.152538
      161    0.133274
      202    0.126395
      230    0.153914
      248    0.127771
      276    0.125019
      287    0.125019
      346    0.127771
      Name: s1, dtype: float64
```

```
[152]: # Upper Outliers for s1
df_diabetes.loc[(df_diabetes['s1']>0.12), 's1']
```

```
[152]: 123    0.152538
      161    0.133274
      202    0.126395
      230    0.153914
      248    0.127771
      276    0.125019
      287    0.125019
      346    0.127771
      Name: s1, dtype: float64
```

```
[153]: # Upper Outliers for s3
df_diabetes.loc[(df_diabetes['s3'] > 0.12), 's3']
```

```
[153]: 35      0.133318
      58      0.181179
      260     0.151726
      261     0.177497
      266     0.122273
      269     0.159089
      286     0.140681
      433     0.122273
      441     0.173816
      Name: s3, dtype: float64
```

```
[154]: # Upper Outliers for s6
      df_diabetes.loc[(df_diabetes['s6'] < -0.125), 's6']
```

```
[154]: 84      -0.129483
      245     -0.129483
      406     -0.137767
      Name: s6, dtype: float64
```

```
[155]: # Lower Outliers for s6
      df_diabetes.loc[(df_diabetes['s6'] < -0.12), 's6']
```

```
[155]: 84      -0.129483
      245     -0.129483
      406     -0.137767
      Name: s6, dtype: float64
```

1.5 Method 4: Z-score using Python

z-score is measure of how many standard deviations a value is from the mean.


Per investopedia.com, “a z-score can be used to determine how far a stock’s return differs from it’s average return. Z-scores are measures of an instrument’s variability and can be used by traders to help determine volatility.”

```
[157]: import statsmodels.api as sm
```

Visual inspection shows most of the values for each column look normally distributed. A normal distribution is a requirement for z-scores. Histograms can also highlight outliers. Look at the chart for s4. The right-side shows outliers.

```
[190]: plt.figure(figsize=(25,15))
      plt.rc('font', size=15)

      for indx, colnm in enumerate(cols):
          plt.subplot(3,4, indx+1)
          sns.histplot(df_diabetes[colnm], color='orange')
```



1.5.1 Calculate z-scores for each column

```
[161]: import scipy.stats as stats
df_zscore = df_diabetes.select_dtypes(include='number').apply(stats.zscore)
df_zscore = df_zscore.drop(['age', 'sex', 'target'], axis=1)
df_zscore
```

```
[161]:
```

| | bmi | bp | s1 | s2 | s3 | s4 | s5 \ |
|-----|-----------|-----------|-----------|-----------|-----------|-----------|-----------|
| 0 | 1.297088 | 0.459841 | -0.929746 | -0.732065 | -0.912451 | -0.054499 | 0.418531 |
| 1 | -1.082180 | -0.553505 | -0.177624 | -0.402886 | 1.564414 | -0.830301 | -1.436589 |
| 2 | 0.934533 | -0.119214 | -0.958674 | -0.718897 | -0.680245 | -0.054499 | 0.060156 |
| 3 | -0.243771 | -0.770650 | 0.256292 | 0.525397 | -0.757647 | 0.721302 | 0.476983 |
| 4 | -0.764944 | 0.459841 | 0.082726 | 0.327890 | 0.171178 | -0.054499 | -0.672502 |
| .. | ... | ... | ... | ... | ... | ... | ... |
| 437 | 0.413360 | 1.256040 | -0.119769 | -0.053957 | -0.602843 | -0.054499 | 0.655787 |
| 438 | -0.334410 | -1.422086 | 1.037341 | 1.664355 | -0.602843 | 0.721302 | -0.380819 |
| 439 | -0.334410 | 0.363573 | -0.785107 | -0.290965 | -0.525441 | -0.232934 | -0.985649 |
| 440 | 0.821235 | 0.025550 | 0.343075 | 0.321306 | -0.602843 | 0.558384 | 0.936163 |
| 441 | -1.535374 | -1.711613 | 1.760535 | 0.584649 | 3.654268 | -0.830301 | -0.088752 |
| | s6 | | | | | | |
| 0 | -0.370989 | | | | | | |
| 1 | -1.938479 | | | | | | |
| 2 | -0.545154 | | | | | | |
| 3 | -0.196823 | | | | | | |


```

4    -0.980568
..    ...
437    0.151508
438    0.935254
439    0.325674
440   -0.545154
441    0.064426

```

```
[442 rows x 8 columns]
```

```
[162]: # Outliers based on z-score. Three standard deviations from the mean.
df_zscore.loc[(df_zscore['bmi'] > 2.576) | (df_zscore['bmi'] < -2.576), 'bmi' ]
```

```

[162]: 32      2.634011
      145      2.701990
      256      3.381781
      262      2.679330
      366      2.883268
      367      3.585718
      405      2.588691
      Name: bmi, dtype: float64

```

```
[163]: # Outliers based on z-score. Three standard deviations from the mean.
df_zscore.loc[(df_zscore['bmi'] > 2.576) | (df_zscore['s1'] < -2.576), 's1' ]
```

```

[163]: 32      -1.132240
      76      -2.665411
      145     -0.698324
      256     -0.611541
      262      0.343075
      366      0.863775
      367      0.632353
      405     -2.202567
      Name: s1, dtype: float64

```

1.6 Method 4: z-score using SQL

1.6.1 Finding Outliers in an OracleLive SQL database using z-scores.

Assumption: Data has a normal distribution

CAL_MONTH_SALES_MV does not have any outliers

```

2  -- Oracle dummy tables.
3  -- Calculate z-score
4  -- Outliers: data that falls outside 3 standard deviations from the mean
5  -- CAL_MONTH_SALES_MV does not have any outliers
6
7  select *
8  from (
9    select calendar_month_desc, dollars
10   , round((dollars - (avg(dollars) over() ))/stddev(dollars) over(),2) as zscore
11   from SH.CAL_MONTH_SALES_MV
12   ) score_table
13   where zscore > 2.576 or zscore < -2.576;
14

```

no data found

Assumption: Data has a normal distribution

CO.STORE_ORDERS has two outliers.

```

21
22  -- Oracle dummy tables.
23  -- Calculate z-score.
24  -- Assumption: data is normally distributed
25  -- Outliers: data that falls outside 3 standard deviations from the mean
26  -- CO.STORE_ORDERS has two outliers
27
28  select *
29  from (
30    select store_name, total_sales
31   , round((total_sales - (avg(total_sales) over() ))/stddev(total_sales) over(),2) as zscore
32   from co.store_orders
33   where order_status = 'COMPLETE'
34   ) score_table
35   where zscore > 2.576 or zscore < -2.576;

```

| STORE_NAME | TOTAL_SALES | ZSCORE |
|------------|-------------|--------|
| Online | 211107.78 | 2.58 |
| - | 299889.62 | 3.81 |

[]: